R and Power BI Project

Hollywood's Most Profitable Stories dataset

Alena Pavlioglo

Just IT Data Analyst Bootcamp

23/03/2023

The goal of this lab is to learn the statistical concepts and to analyse “Hollywood's Most Profitable Stories”. dataset on Power BI.

CONTENT

Data Preparation

Load data and view the data.

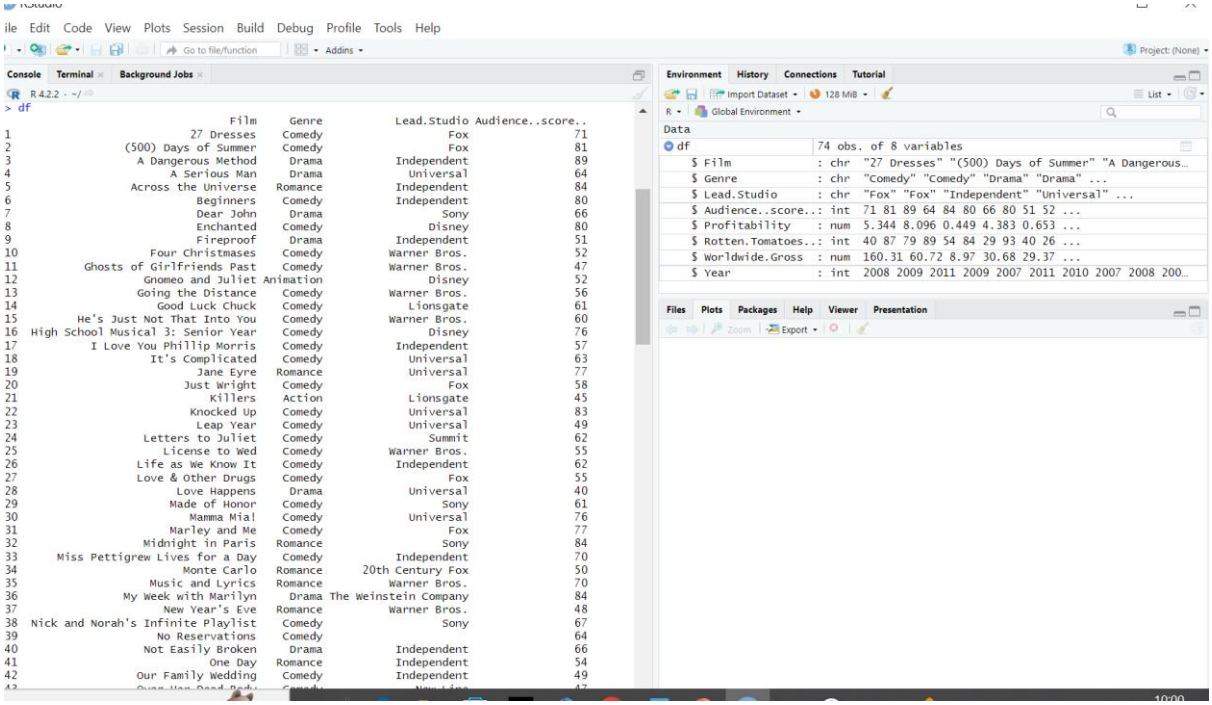

To view the data

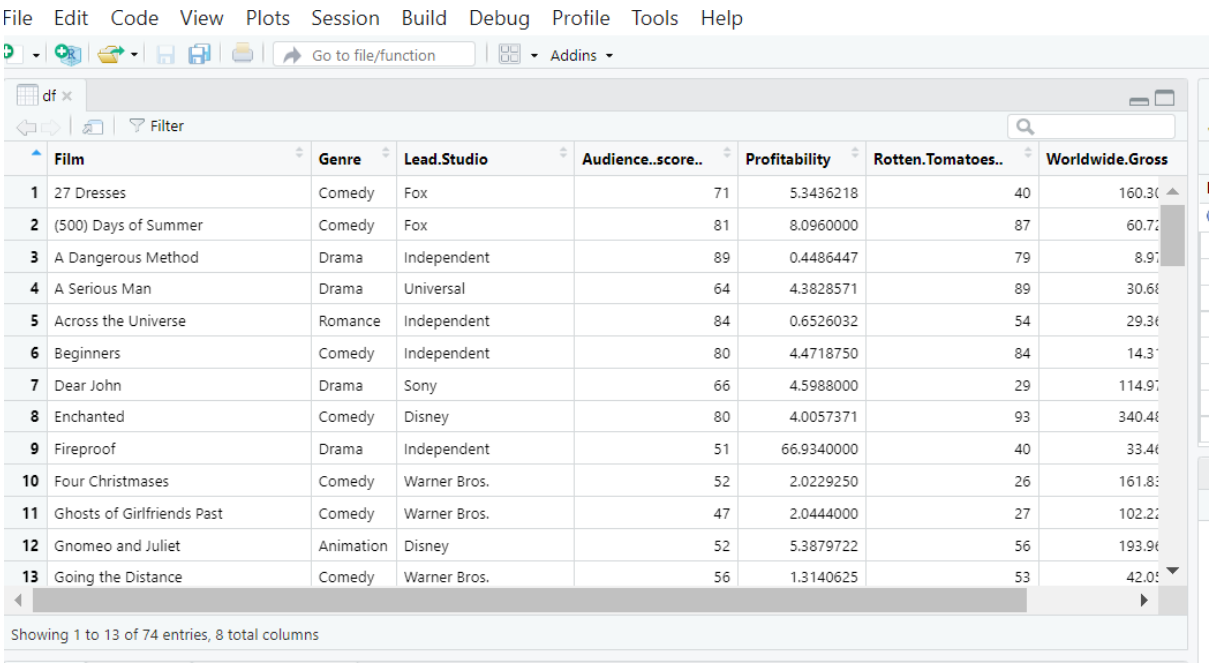

df

Filter

	Film	Genre	Lead.Studio	Audience..score..	Profitability	Rotten.Tomatoes..	Worldwide.Gross
1	27 Dresses	Comedy	Fox	71	5.3436218	40	160.3(
2	(500) Days of Summer	Comedy	Fox	81	8.0960000	87	60.7
3	A Dangerous Method	Drama	Independent	89	0.4486447	79	8.9
4	A Serious Man	Drama	Universal	64	4.3828571	89	30.6
5	Across the Universe	Romance	Independent	84	0.6526032	54	29.3
6	Beginners	Comedy	Independent	80	4.4718750	84	14.3
7	Dear John	Drama	Sony	66	4.5988000	29	114.9
8	Enchanted	Comedy	Disney	80	4.0057371	93	340.4
9	Fireproof	Drama	Independent	51	66.9340000	40	33.4
10	Four Christmases	Comedy	Warner Bros.	52	2.0229250	26	161.8
11	Ghosts of Girlfriends Past	Comedy	Warner Bros.	47	2.0444000	27	102.2
12	Gnomeo and Juliet	Animation	Disney	52	5.3879722	56	193.9
13	Going the Distance	Comedy	Warner Bros.	56	1.3140625	53	42.0

Showing 1 to 13 of 74 entries, 8 total columns

The head() function in R is used to display the first *n* rows present in the input data frame.

Head(df)

```
> head(df)
                Film   Genre Lead.Studio Audience..score.. Profitability Rotten.Tomatoes..
1          27 Dresses  Comedy         Fox                71     5.3436218                40
2 (500) Days of Summer  Comedy         Fox                81     8.0960000                87
3   A Dangerous Method   Drama Independent                89     0.4486447                79
4        A Serious Man   Drama   Universal                64     4.3828571                89
5  Across the Universe Romance Independent                84     0.6526032                54
6            Beginners  Comedy Independent                80     4.4718750                84
  Worldwide.Gross Year
1      160.308654 2008
2       60.720000 2009
3        8.972895 2011
4       30.680000 2009
5       29.367143 2007
6       14.310000 2011
> |
```

By checking the column names and structure I found out that there are 8 variables and 74 objects.

```
> colnames(df)
[1] "Film"              "Genre"             "Lead.Studio"       "Audience..score.."
[5] "Profitability"     "Rotten.Tomatoes.." "Worldwide.Gross"   "Year"
> str(df)
'data.frame':   74 obs. of  8 variables:
 $ Film             : chr  "27 Dresses" "(500) Days of Summer" "A Dangerous Method" "A Serious Man" ...
 $ Genre            : chr  "Comedy" "Comedy" "Drama" "Drama" ...
 $ Lead.Studio      : chr  "Fox" "Fox" "Independent" "Universal" ...
 $ Audience..score..: int  71 81 89 64 84 80 66 80 51 52 ...
 $ Profitability    : num  5.344 8.096 0.449 4.383 0.653 ...
 $ Rotten.Tomatoes..: int  40 87 79 89 54 84 29 93 40 26 ...
 $ Worldwide.Gross  : num  160.31 60.72 8.97 30.68 29.37 ...
 $ Year             : int  2008 2009 2011 2009 2007 2011 2010 2007 2008 2008 ...
> |
```

Load packages

To be able to work and use different functions to make statistical and graphical analysis on datasets we have to load some libraries.

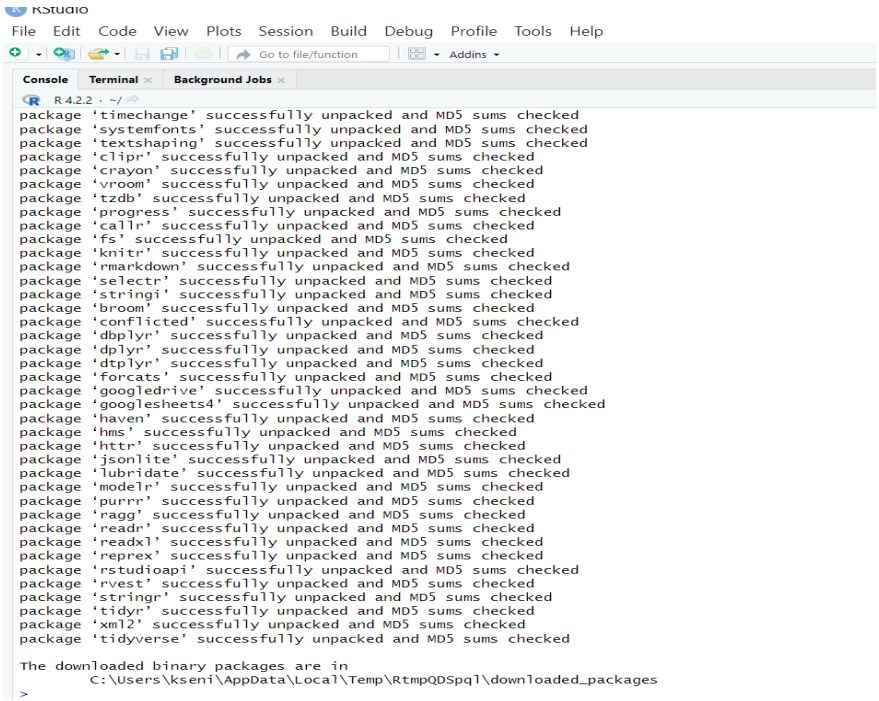

Check data types:

```
> str(df)
'data.frame':   74 obs. of  8 variables:
 $ Film             : chr  "27 Dresses" "(500) Days of Summer" "A Dangerous Method" "A Serious Man" ...
 $ Genre            : chr  "Comedy" "Comedy" "Drama" "Drama" ...
 $ Lead.Studio      : chr  "Fox" "Fox" "Independent" "Universal" ...
 $ Audience..score..: int  71 81 89 64 84 80 66 80 51 52 ...
 $ Profitability    : num  5.344 8.096 0.449 4.383 0.653 ...
 $ Rotten.Tomatoes..: int  40 87 79 89 54 84 29 93 40 26 ...
 $ Worldwide.Gross  : num  160.31 60.72 8.97 30.68 29.37 ...
 $ Year             : int  2008 2009 2011 2009 2007 2011 2010 2007 2008 2008 ...
```

To access the data in a single column to explore ethe data (for example column Genre)

```
> df$Genre
 [1] "Comedy"    "Comedy"    "Drama"     "Drama"     "Romance"   "Comedy"    "Drama"
 [8] "Comedy"    "Drama"     "Comedy"    "Comedy"    "Animation" "Comedy"    "Comedy"
[15] "Comedy"    "Comedy"    "Comedy"    "Comedy"    "Romance"   "Comedy"    "Action"
[22] "Comedy"    "Comedy"    "Comedy"    "Comedy"    "Comedy"    "Comedy"    "Drama"
[29] "Comedy"    "Comedy"    "Comedy"    "Romance"   "Comedy"    "Romance"   "Romance"
[36] "Drama"     "Romance"   "Comedy"    "Comedy"    "Drama"     "Romance"   "Comedy"
[43] "Comedy"    "Romance"   "Comedy"    "Drama"     "Drama"     "Comedy"    "Comedy"
[50] "Comedy"    "Romance"   "Animation" "Comedy"    "Fantasy"   "Drama"     "Comedy"
[57] "Comedy"    "Comedy"    "Drama"     "Drama"     "Comedy"    "Romance"   "Romance"
[64] "Romance"   "Comedy"    "Romance"   "Romance"   "Animation" "Drama"     "Comedy"
[71] "Comedy"    "Comedy"    "Comedy"    "Romance"
>
```

1. Data Cleaning

The purpose of data cleaning is to identify, correct, or remove inaccurate raw data for downstream purposes.

Checking missing values. To find the length of columns for missing values we use colSum(is.na(df)) and to remove all rows that contains at least on NA we use command df <- na.omit(df) and check if the rows have been removed. It is very important to handle missing values since it can bias the results and reduce the accuracy.

```
> colSums(is.na(df))
             Film             Genre       Lead.Studio Audience..score..
                0                 0                 0                 1
    Profitability Rotten.Tomatoes..   Worldwide.Gross              Year
                3                 1                 0                 0
>
> df <- na.omit(df)
>
> colSums(is.na(df))
             Film             Genre       Lead.Studio Audience..score..
                0                 0                 0                 0
    Profitability Rotten.Tomatoes..   Worldwide.Gross              Year
                0                 0                 0                 0
>
```

Removing duplicates and rounding the values to 2 places. Checking the dimension of the new data frame.
Removing duplicates from data set is also very important in order to maintain accuracy and avoid misleading statistics. Then we specify the number of decimal places (2) to which we need to round the “Profitability” and “Worldwide.Gross” Columns.

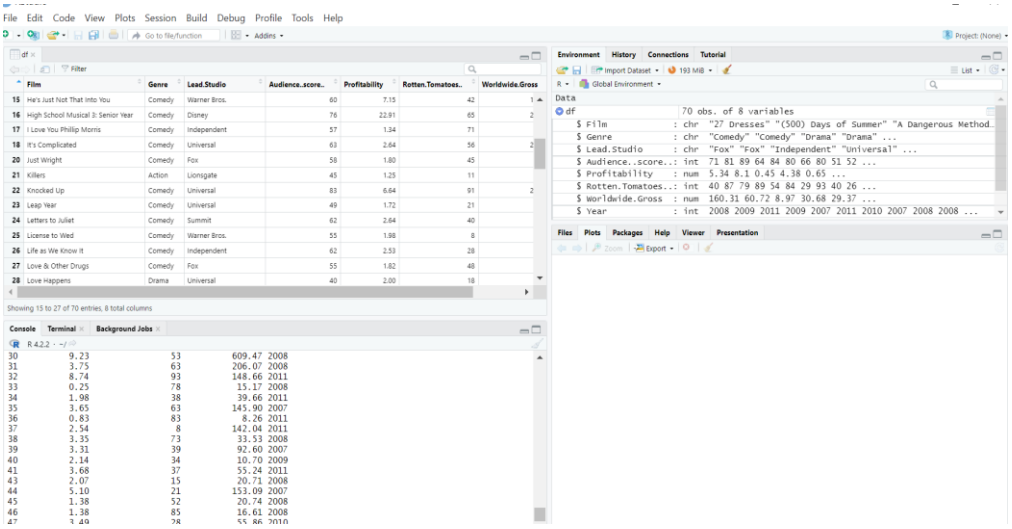

Step 2.1

Check for outliers using a boxplot.

Boxplot is the useful tool to detect potential outliers and helps to visualize a quantitative variable by displaying minimum, median, first and third quartiles and maximum and any observation that was classified as a suspected outlier using the IQR criterion. We create a boxplot and to adjust the y-axis we use coord_cartesian and set up y-axis range from 0 to 1000, just as we specified using ylim() argument. To label the scale of the x-axis we have applied scale_x_continuous to change x-axis when x-variable is continuous.

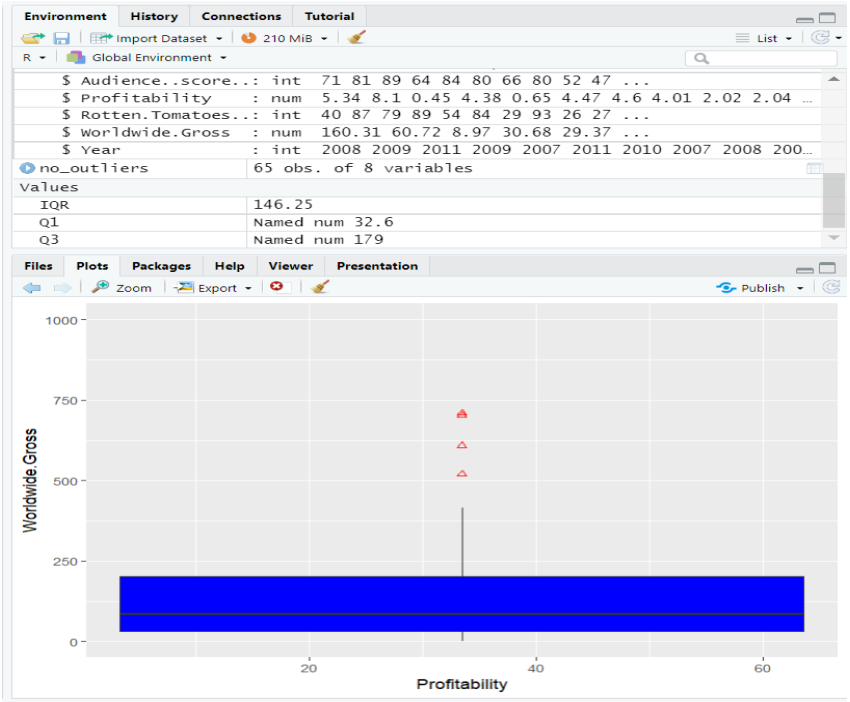

Quartiles are three values that split dataset into quarters.
Q1 First quartile: 25% of all the values fall below that value.
Q2: Second quartile / Median: This value splits the data in half.
Q3 Third quartile: 25% of the data are above this value.

The interquartile range (IQR) is the range between the first and third quartiles. IQR = Q3 -Q1

Observations considered as potential outliers by the IQR criterion are displayed as points in the boxplot.
By removing outliers in" Profitability" and "Worldwide.Gross" and put a condition in which new subset will meet criteria that all observation above
Q3 +1.5*IQR and below Q1 -1.5*IQR are considered as potential outliers and will be removed.

```
> Q1 <- quantile(df$Profitability, .25)
>
> Q3 <- quantile(df$Profitability, .75)
>
> IQR <- IQR(df$Profitability)
>
> no_outliers <- subset(df, df$Profitability> (Q1 - 1.5*IQR) & df$Profitability< (Q3 + 1.5*IQR))
>
> dim(no_outliers)
[1] 65  8
>
>
```

The dimensions for new "no_outliers" are 65 objects and 8 attributes, but after removing the outliers for "Worldwide.Gross" we notice that we have a new df1 with 61 objects and 8 attributes.

```
> Q1 <- quantile(no_outliers$Worldwide.Gross, .25)
>
> Q3 <- quantile(no_outliers$Worldwide.Gross, .75)
>
> IQR <- IQR(no_outliers$Worldwide.Gross)
> df1 <- subset(no_outliers, no_outliers$Worldwide.Gross> (Q1 - 1.5*IQR) & no_outliers$Worldwide.Gross< (Q3 +
1.5*IQR))
> dim(df1)
[1] 61  8
```

There are various methods for extracting the values of the potential outliers and I was wondering which of them are more accurate, how they work for R and whether I will come to the same results.
One possible way also based on the IQR criterion by using boxplot.stats()$out and function which()
to extract the row number corresponding to these outliers. We can also try z-score and Hampel filter by using median() and mad() functions.

```
> df$z_Profitability<- scale(df$Profitability)
> hist(df$z_Profitability)
> summary(df$z_Profitability)
       V1
 Min.   :-0.57342
 1st Qu.:-0.35708
 Median :-0.25640
 Mean   : 0.00000
 3rd Qu.: 0.02286
 Max.   : 7.44847
> which(df$z_Profitability > 3.29)
[1] 9
> lower_bound <- median(df$Profitability) - 3 * mad(df$Profitability, constant = 1)
> lower_bound
[1] -1.15
> upper_bound <- median(df$Profitability) + 3 * mad(df$Profitability, constant = 1)
> upper_bound
[1] 6.44
> outlier_ind <- which(df$Profitability < lower_bound | df$Profitability > upper_bound)
> outlier_ind
 [1]  2  9 15 16 21 29 31 46 55 57 59 64
> boxplot.stats(df$Profitability)$out
[1] 66.93 22.91 14.20 10.18 11.09
> out_ind <- which(df$Profitability %in% c(out))
> out_ind
[1]  9 16 57 59 64
>
```

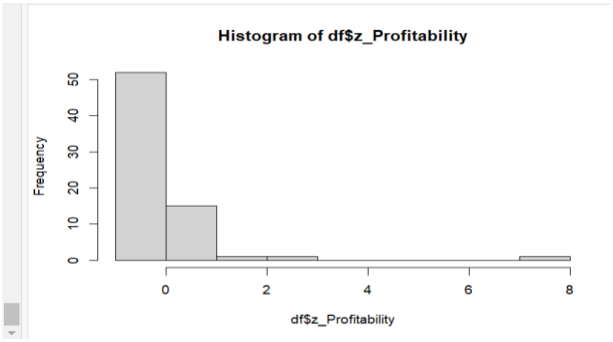

Step 3: Exploratory Data Analysis

The summary is an exploratory data analysis tool that provides insight into the distribution of values for one Variable. This set of statistics describes where data values occur, their central tendency, variability, and the general shape of their distribution.

```
> summary(df1)
     Film               Genre           Lead.Studio        Audience..score.. Profitability
 Length:61          Length:61          Length:61          Min.   :35.00     Min.   :0.000
 Class :character   Class :character   Class :character   1st Qu.:52.00     1st Qu.:1.750
 Mode  :character   Mode  :character   Mode  :character   Median :62.00     Median :2.530
                                                          Mean   :63.02     Mean   :3.014
                                                          3rd Qu.:72.00     3rd Qu.:3.750
                                                          Max.   :89.00     Max.   :8.740

 Rotten.Tomatoes.. Worldwide.Gross       Year
 Min.   : 3.0      Min.   :  0.03   Min.   :2007
 1st Qu.:27.0      1st Qu.: 32.40   1st Qu.:2008
 Median :43.0      Median : 69.31   Median :2009
 Mean   :46.7      Mean   :103.16   Mean   :2009
 3rd Qu.:64.0      3rd Qu.:153.09   3rd Qu.:2010
 Max.   :93.0      Max.   :355.08   Max.   :2011
> |
```

From the summary we can see that the Median for “Worldwide.Gross”(69.31) is close to Q1(32.40) thanQ3 (153.09). Therefore, the distribution of values is right- skewed.

Bivariate analysis refers to the analysis of two variables to determine relationships between them.

Scatterplot

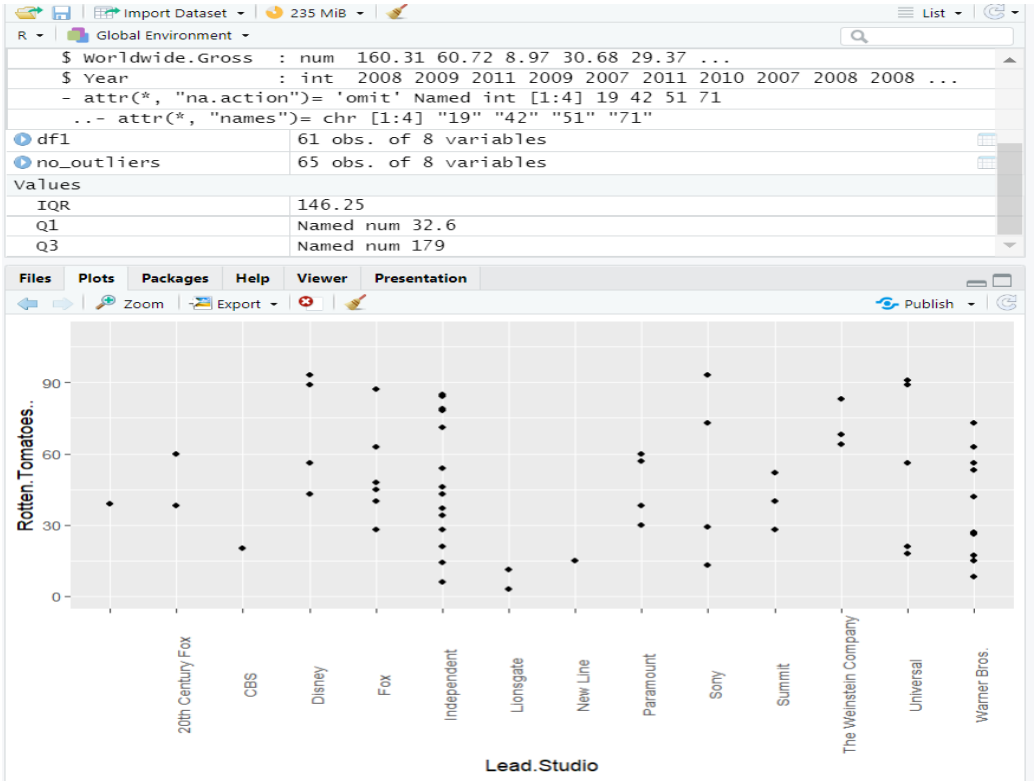

To explore more interesting options for displaying the scatterplots I uploaded a “skatterplot3d” package and build 3d scatterplot in Films, Genres and Lead.Studious.

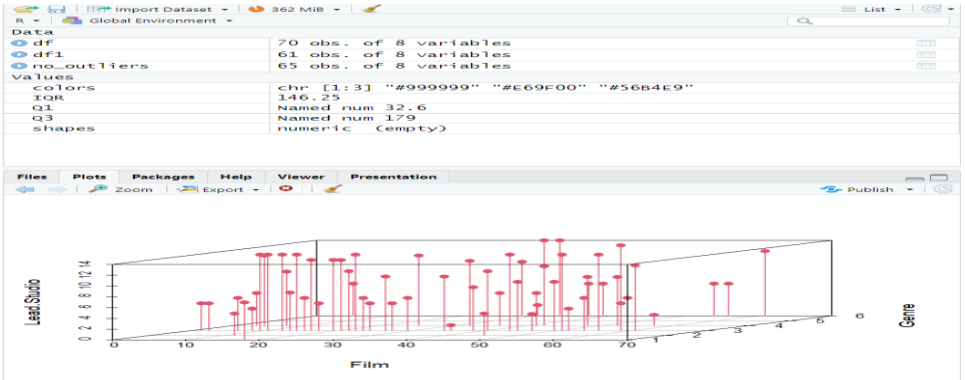

Bar chart

I have created a bar chat by Year and grouped it by Genre.

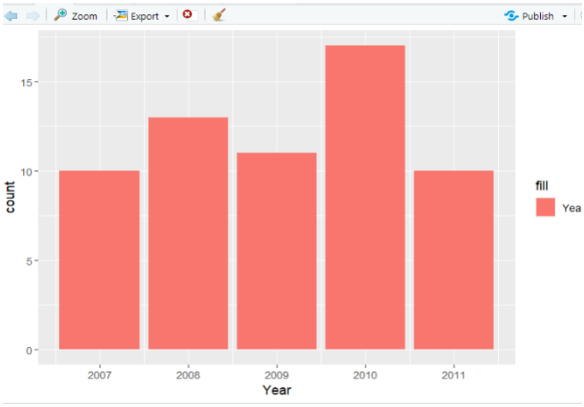

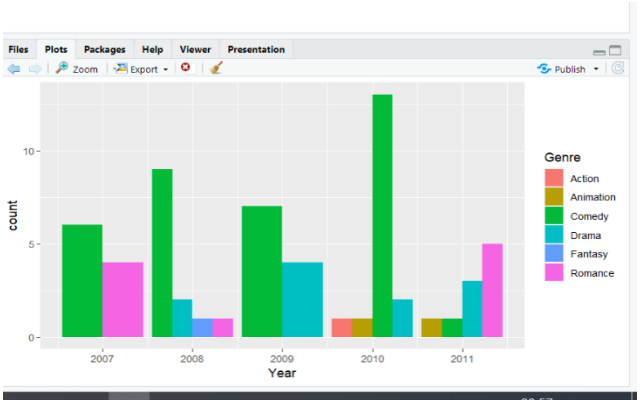

Density plot

I was really interested to build a Density plot since it visualizes the distribution of data over a continuous interval or time period. The peaks of a Density Plot help display where values are concentrated over the interval. I have chosen 2 different parameters and visualized the density by Genre and Profitability aver the Year.

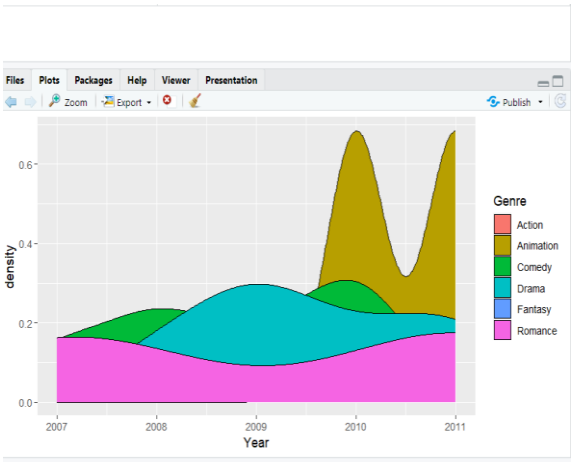

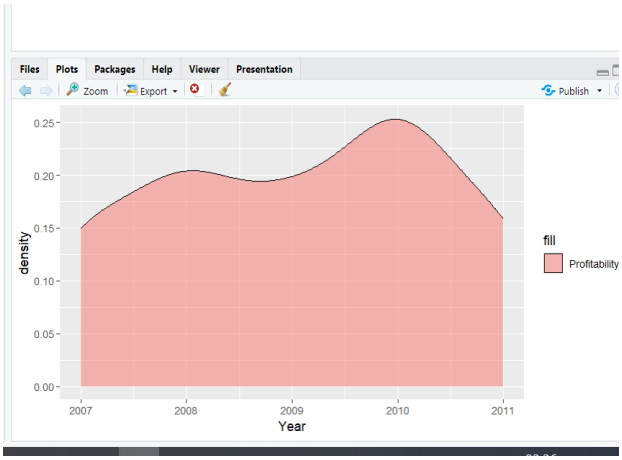

Step 4: Export data

Finally, our clean data is ready do be exported to Power BI to analyze the data and visualize results by creating different charts. We upload our already cleaned dataset as a CSV file to Power BI.

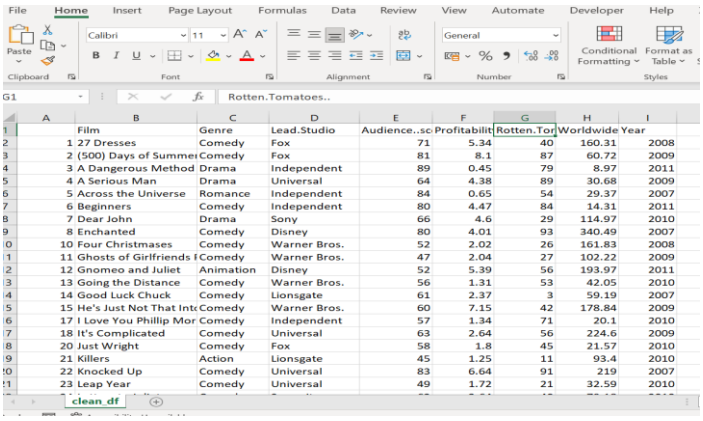

	Film	Genre	Lead.Studio	Audience..sc	Profitabilit	Rotten.Tor	Worldwide	Year
1	27 Dresses	Comedy	Fox	71	5.34	40	160.31	2008
2	(500) Days of Summer	Comedy	Fox	81	8.1	87	60.72	2009
3	A Dangerous Method	Drama	Independent	89	0.45	79	8.97	2011
4	A Serious Man	Drama	Universal	64	4.38	89	30.68	2009
5	Across the Universe	Romance	Independent	84	0.65	54	29.37	2007
6	Beginners	Comedy	Independent	80	4.47	84	14.31	2011
7	Dear John	Drama	Sony	66	4.6	29	114.97	2010
8	Enchanted	Comedy	Disney	80	4.01	93	340.49	2007
10	Four Christmases	Comedy	Warner Bros.	52	2.02	26	161.83	2008
11	Ghosts of Girlfriends I	Comedy	Warner Bros.	47	2.04	27	102.22	2009
12	Gnomeo and Juliet	Animation	Disney	52	5.39	56	193.97	2011
13	Going the Distance	Comedy	Warner Bros.	56	1.31	53	42.05	2010
14	Good Luck Chuck	Comedy	Lionsgate	61	2.37	3	59.19	2007
15	He's Just Not That Int	Comedy	Warner Bros.	60	7.15	42	178.84	2009
17	I Love You Phillip Mor	Comedy	Independent	57	1.34	71	20.1	2010
18	It's Complicated	Comedy	Universal	63	2.64	56	224.6	2009
20	Just Wright	Comedy	Fox	58	1.8	45	21.57	2010
21	Killers	Action	Lionsgate	45	1.25	11	93.4	2010
22	Knocked Up	Comedy	Universal	83	6.64	91	219	2007
23	Leap Year	Comedy	Universal	49	1.72	21	32.59	2010

2. The average Rotten Tomatoes ratings of each genre

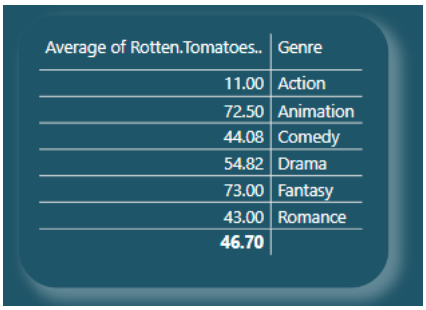

Average of Rotten.Tomatoes..	Genre
11.00	Action
72.50	Animation
44.08	Comedy
54.82	Drama
73.00	Fantasy
43.00	Romance
46.70	

3. The number of movies produced per year.

I created a treemap to visualize the number of movies produced per Year.

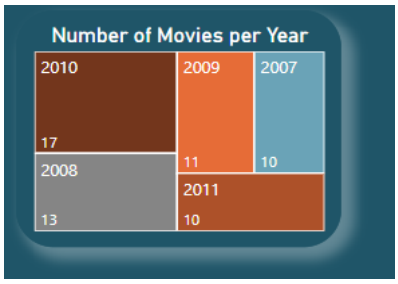

4. The audience scores for each film.

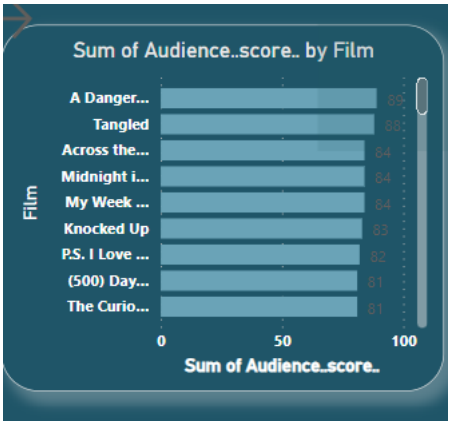

Clustered bar Chart seems to be suitable for analysing this relationship since we have a large amount of data that can be grouped.

5. The profitability per studio.

I have created stacked column chart displays the contribution of each Studio's profitability per Year by using geom_bar() and position = "dodge" to make it "side by side". I have use scale_fill_manual() to specify the colours and labelled the titles.

I wanted to add an extra touch to my bar charts, so I added a line representing an average

of all the bars. In my example, this would give us an insight into which Studio over which Year performed better than average.

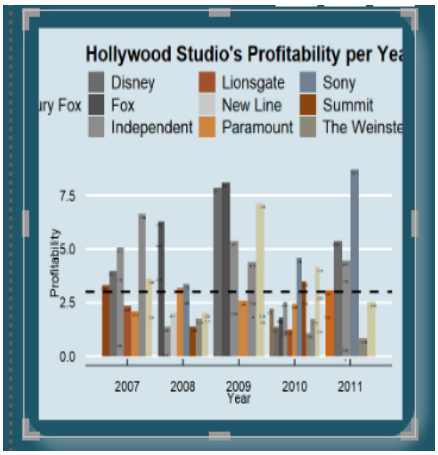

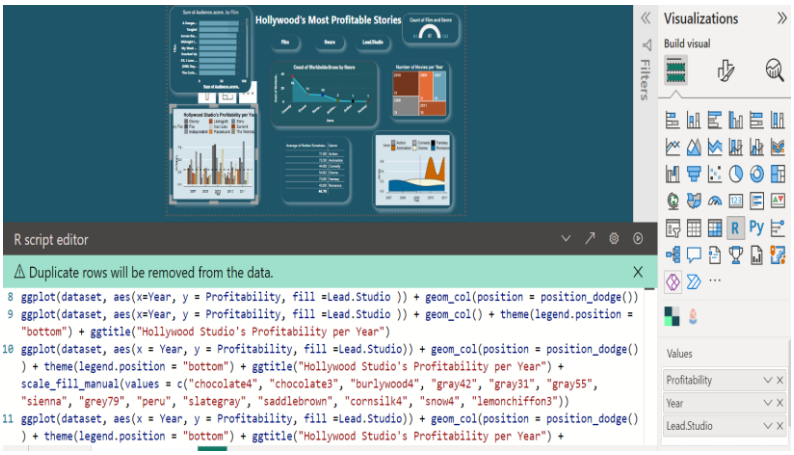

6. The worldwide gross per genre.

I used area chart to visualize and we can see that Comedy Genre has the highest Worldwide Gross and it is the leader, since the rest of the Genres have significant difference.

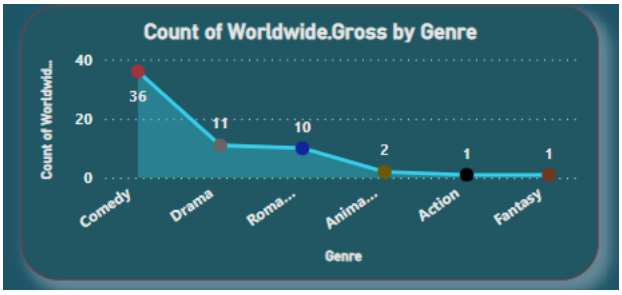

I decided to create a density plot for Genre per Year in R by using geom_density() function and applied theme_economist() + scale_color_economist() to change the background colour.

I also use the geom_vline() layer to add a vertical line for the mean of Year

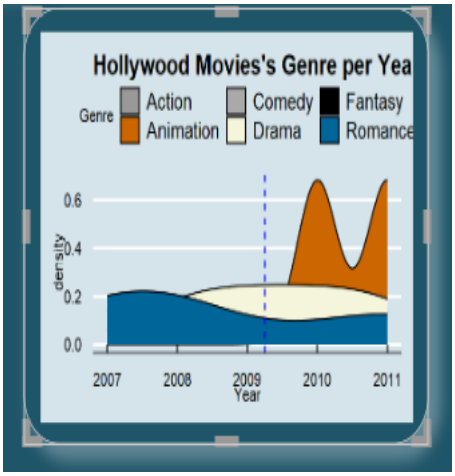

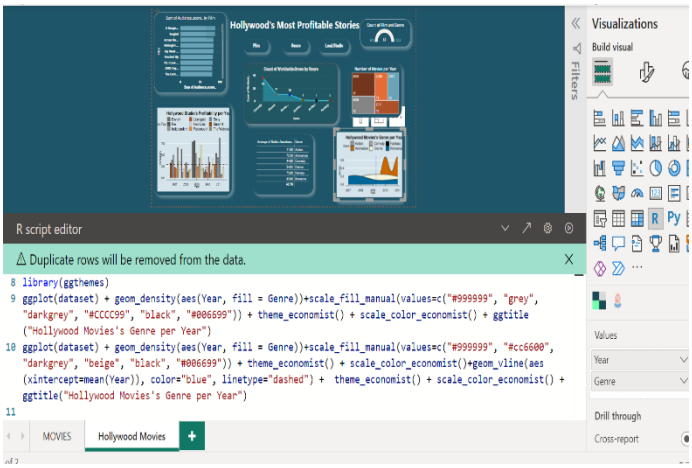

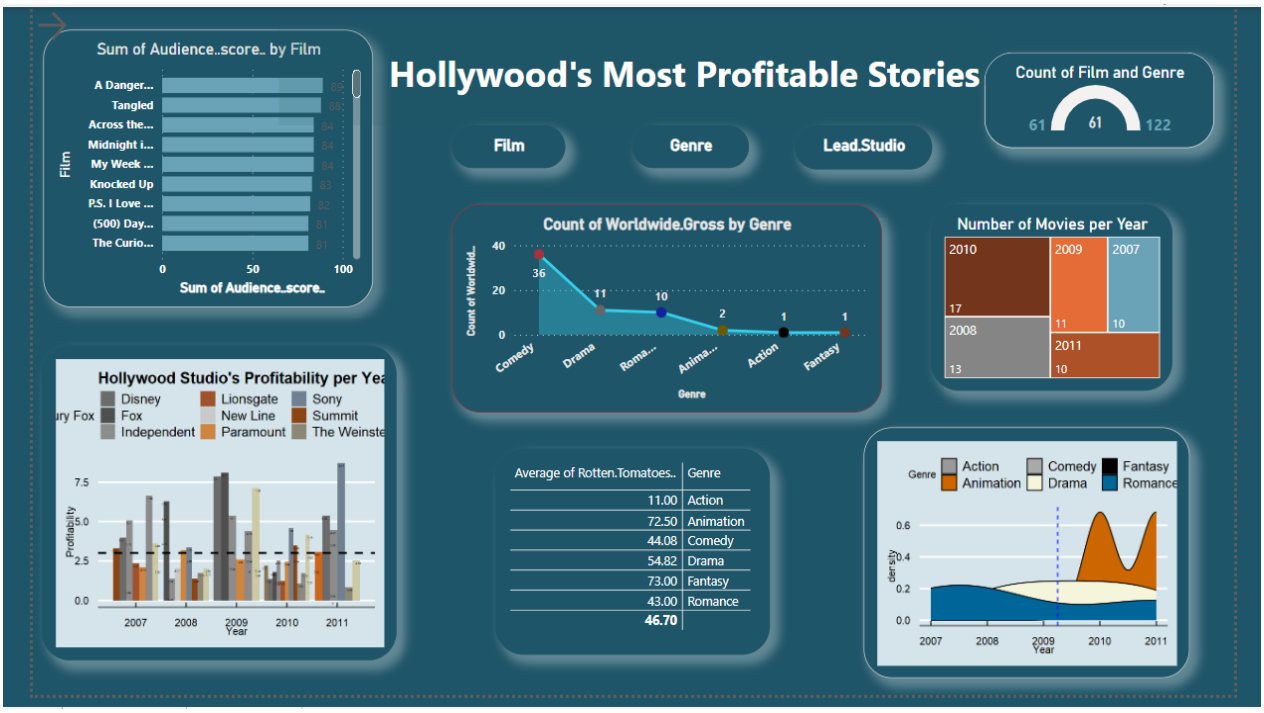

https://app.powerbi.com/links/PtZ6SXLsJ3?ctid=6efd0f20-57c8-4447-b53f-00d4992ca50b&pbi_source=linkShare